GALGOTIAS UNIVERSITY — SPARKX 3.0 INTERNATIONAL

LabSPOC AI: Spatially Grounded Virtual Lab Ambassador & Institutional Intelligence Platform

Final Project Technical Report & Architectural Specification

Challenge Track:	Smart Campus & Digital Twin / AI-Powered Intelligent Lab SPOC
Host Institution:	Oriental Institute of Science & Technology (OIST), Bhopal (M.P.), India
Department & Facility:	Department of Computer Science & Engineering · AI & ML Laboratory (Room 302)
Project Lead / Author:	Shivang Saxena (Individual Innovation Challenge Entry)
Contact Email:	shivangsaxena1011@gmail.com
Preferred AI Provider:	Google Gemini 2.5 Flash & Gemini Embedding-001 (3072 Dimensions)
GitHub Repository:	https://github.com/shivangsaxena1011/LabSPOC-AI-SparkX-3.0
Release Version:	Release 0.9.0 (SPARKX 3.0 Final Submission Package)
Date of Evaluation:	November 10, 2026

Institutional Abstract:
LabSPOC AI introduces a spatially grounded virtual laboratory ambassador and institutional intelligence platform designed for the AI & ML Laboratory at Oriental Institute of Science & Technology (OIST), Bhopal. The system addresses critical knowledge silos in modern university research infrastructure by uniting an interactive 3D digital twin, a two-stage hybrid semantic/lexical RAG engine, an automated Evidence Gate enforcing honest abstention, a persona-tailored Agenda-to-Tour Compiler, and a Spatially Anchored Evidence Graph (SAEG). Backed by 40 automated passing test cases and benchmarked retrieval metrics (100% Recall@5, 1.0000 MRR), the platform establishes an auditable, hallucination-resistant, and privacy-compliant paradigm for smart campus facilities.

Candidate Declaration & Certificate of Authorship

I, **Shivang Saxena**, student and researcher at the Oriental Institute of Science & Technology (OIST), Bhopal, hereby solemnly declare that the project titled **'LabSPOC AI — Spatially Grounded Virtual Lab Ambassador'**, submitted for evaluation in the *SPARKX 3.0 International 30-Day Innovation Challenge* organized by Galgotias University, represents my original individual architectural and software engineering work. The codebase, documentation, test suites, and empirical benchmarks were designed and implemented specifically for this challenge.

I certify that open-source frameworks (Next.js, FastAPI, Three.js, SQLAlchemy) and Google Gemini AI APIs have been integrated in strict accordance with their open licensing terms. No portion of this work involves fabricated test outcomes, simulated RAG responses, or unauthorized proprietary material. All room coordinates are honestly designated as illustrative pending on-site LiDAR scans.

Table of Contents

Section	Title & Core Technical Topics	Page
1	Executive Summary & Institutional Context	3
2	Problem Statement & Academic Laboratory Pain Points	4
3	Systematic Literature Review & Prior Art Gap Analysis	5
4	System Architecture & Tier-by-Tier Specification	6
5	Relational Data Model & Spatially Anchored Schemas	7
6	Module 1: AI Virtual Lab Guide (3D Digital Twin)	8
7	Module 2: Conversational Lab Assistant (Multilingual & Speech)	9
8	Module 3: RAG Knowledge System & Verified Ingestion Engine	10
9	Module 4 & 5: Project Demonstrator & Visitor Analytics	11
10	Signature Innovation A: Spatially Anchored Evidence Graph (SAEG)	12
11	Signature Innovation B: Persona-Based Agenda-to-Tour Compiler	13
12	Signature Innovation C: Evidence Gate & Honest Abstention Protocol	14
13	Visual Interface Verification & Headless Chromium Validation	15
14	Comprehensive Automated Pytest Suite (40/40 Verification)	16
15	Empirical RAG Benchmark & Latency Evaluation Results	17
16	Security Architecture, Threat Modeling & DPDP Act Privacy	18
17	Production Cloud Deployment Topology & Zero-Cost Financial Plan	19
18	Honest Limitations, Future Roadmap & Verified Academic References	20

1. Executive Summary & Institutional Context

Academic artificial intelligence and machine learning laboratories are dynamic, capital-intensive research hubs. At the Oriental Institute of Science & Technology (OIST), Bhopal, the AI & ML Laboratory (Room 302) serves as the campus flagship for accelerated computing, edge artificial intelligence development, and computer vision experimentation. The laboratory houses critical institutional infrastructure: high-density NVIDIA DGX compute nodes, edge development boards (Jetson AGX Orin, Jetson Nano), multi-camera calibration rigs, and high-performance developer workstations.

Despite the sophistication of this physical infrastructure, university facilities across the globe suffer from a persistent operational paradox: **institutional intelligence remains offline, siloed, and ephemeral**. Prospective students, accreditation evaluators (such as NBA and NAAC visiting committees), visiting industry researchers, and undergraduate scholars experience significant friction when attempting to comprehend lab capabilities, verify safety protocols, or inspect ongoing capstone projects. Faculty in-charge personnel are repeatedly pulled away from active research to conduct repetitive facility walkthroughs.

The LabSPOC AI Vision:

LabSPOC AI introduces an embodied, spatially grounded, and evidence-governed virtual laboratory ambassador (Single Point of Contact - SPOC). Unlike generic campus chatbots that rely on ungrounded conversational models prone to hallucination, LabSPOC AI pairs an interactive 3D WebGL digital twin with a dual-stage hybrid retrieval engine (combining dense 3072-dimensional Gemini embeddings with lexical BM25-style term matching) and a deterministic verification layer termed the **Evidence Gate**.

The system serves two distinct operational modalities: a high-efficiency **Operations & Administration Workspace** for lab coordinators to manage equipment, review uploaded documentation, and monitor telemetry; and an engaging, accessible **Visitor Experience Portal** featuring multi-turn conversational intelligence, native bilingual voice interaction (English/Hindi), and persona-driven guided itineraries.

Institutional Dimension	Conventional University Reality	LabSPOC AI Transformed State
Facility Exploration	Physical walk-in required; static photos or un-navigable Matterport pans.	Interactive 3D WebGL twin with responsive hotspots and camera tweening.
Information Grounding	Scattered across paper binders, local PDFs, or faculty memory.	Hybrid RAG with 100% citation verification and deterministic Evidence Gate.
Visitor Alignment	One-size-fits-all verbal tour or unguided wander.	Algorithmic Agenda-to-Tour compilation for 2, 5, or 15 minute budgets.
Research Visibility	Capstone projects archived in private Git repos after semester ends.	Dynamic Project Demonstrator linked to utilized hardware assets.

2. Problem Statement & Academic Laboratory Pain Points

Modern university research laboratories operate in high-friction environments characterized by four critical structural bottlenecks. Through extensive workflow analysis of the AI & ML Laboratory at OIST Bhopal, we formalized these pain points:

2.1 Computational Asset Invisibility & Policy Friction

Advanced computing machinery (e.g., enterprise GPU nodes, specialized sensor rigs) represents significant capital expenditure. However, student researchers often remain unaware of available cluster hardware, Slurm queue priorities, or container prerequisites. Standard Operating Procedures (SOPs), emergency power-down sequences, and high-voltage electrical safety guidelines exist in physical binders that are rarely consulted until an equipment failure or safety incident occurs.

2.2 Administrative Overhead in Stakeholder Walkthroughs

Each semester, academic laboratories host dozens of disparate delegations: high-level industry recruiters, government accreditation teams, international researchers, and prospective undergraduate scholars. Each cohort possesses differing technical depth and time constraints. A corporate recruiter seeking compute partnership needs a 5-minute executive briefing; an accreditation auditor requires verified asset logs; a prospective researcher wants deep project demonstrations. Providing tailored walkthroughs manually consumes hundreds of faculty hours.

2.3 The Critical Risk of Conversational LLM Hallucinations

When institutions deploy conversational AI chatbots to automate visitor queries, they encounter the severe hazard of generative hallucinations. Standard LLMs generate plausible-sounding but completely fabricated specifications (e.g., falsely asserting an edge cluster supports CUDA 12.4 when running a legacy JetPack build, or inventing fictitious lab safety rules). In an active laboratory environment with energized high-power equipment, false AI claims introduce unacceptable physical liability and institutional reputational risk.

2.4 Capstone Research Isolation & Project Sinks

Undergraduate and postgraduate students build remarkable machine learning systems—such as automated agricultural disease detection, low-latency edge robotics, and biomedical imaging pipelines. Once graded, these capstones typically disappear into archived folders. Visiting VIPs and corporate sponsors leave the facility without discovering the innovative research developed right inside the room.

Pain Point	Root Cause	Operational Impact	LabSPOC AI Countermeasure
Asset Invisibility	Siloed physical manuals & disconnected asset lists.	Under-utilized compute; safety guideline ignorance.	SAEG: Spatially Anchored Evidence Graph linking docs to 3D beacons.
Tour Coordination	Manual faculty scheduling & inflexible itineraries.	Faculty fatigue; rushed or irrelevant visitor tours.	Agenda-to-Tour Compiler with dynamic time budgeting (2/5/15 min).
AI Hallucinations	Ungrounded text generators without fact checking.	Safety hazards; incorrect technical guidance to visitors.	Evidence Gate enforcing strict citation validity & honest abstention.
Project Sinks	Lack of unified exhibition infrastructure.	Missed hiring, funding, and publication opportunities.	Interactive Project Demonstrator linked to utilized hardware assets.

3. Systematic Literature Review & Prior Art Gap Analysis

To position LabSPOC AI within the academic body of knowledge, we conducted a systematic literature review across three intersecting domains: Virtual Campus & Digital Twin Environments, Retrieval-Augmented Generation (RAG) Architectures, and Embodied Conversational Agents.

3.1 Digital Twins & Spatial Campus Representations

Digital twins in educational institutions historically evolved from 360-degree panoramic photography (e.g. Matterport virtual tours) and Building Information Modeling (BIM) frameworks. While Matterport capture offers high photorealism, it remains fundamentally passive: the environment cannot answer semantic questions, link live research projects to equipment, or adapt to visitor personas. In contrast, WebGL-based scientific visualization engines such as Mol* (Sehnal et al., 2021) demonstrated that browser-native rendering enables rich user interaction without hefty software downloads.

3.2 Retrieval-Augmented Generation & Hallucination Suppression

The seminal work of Lewis et al. (2020) and Guu et al. (2020) established RAG as the preeminent architecture for grounding neural models in external parametric knowledge. However, conventional RAG systems operate in textual vacuums: text passages lack physical coordinates, zone associations, or hardware status. Recent evaluations by Gao et al. (2023) show that naive RAG fails in technical domains when document chunks share ambiguous terminology. LabSPOC AI overcomes this through a two-stage hybrid retrieval strategy combining 3072-dimensional Gemini vector embeddings with exact lexical term matching, further boosted by the user's active spatial hotspot.

3.3 Frontier Multimodal LLMs in Enterprise Workflows

Modern frontier models, including the Gemini 1.5/2.5 series (Gemini Team, 2024), exhibit remarkable multilingual understanding and reasoning capabilities. However, deploying them in public educational facilities requires strict governance: responses must be bounded by verified evidence, citations must refer to authentic source paragraphs, and the system must honestly abstain when data is missing.

System / Approach	3D Spatial Interaction	Conversational Grounding	Hallucination Prevention	Persona Compilation
Standard Virtual Tour (Matterport)	Panoramic 360 images (non-semantic)	None (Static text tags)	N/A (No AI generation)	None (Linear path)
Standard Campus Chatbot	None (2D Chat bubble)	Keyword lookup / Naive LLM	Poor (Frequent hallucinations)	None (Generic script)
Academic RAG Pipeline (Lewis et al.)	None (Text-only)	Dense vector passage retrieval	Moderate (Soft prompt grounding)	None (Single prompt)
LabSPOC AI (This Work)	**Interactive Three.js 3D Twin**	**Dual Hybrid RAG + Spatial Boost**	**Deterministic Evidence Gate**	**Agenda-to-Tour Compiler**

4. System Architecture & Tier-by-Tier Specification

LabSPOC AI is architected as an enterprise-grade, loosely coupled multi-tier distributed platform. The system enforces strict separation of concerns across the Presentation, API Gateway, Intelligence, and Persistence tiers:

[PRESENTATION TIER] Next.js 14 · React 18 · Three.js · Tailwind CSS
- Interactive 3D Digital Twin (`Canvas`, OrbitControls, Custom Shaders, Hotspot Beacons)
- Visitor Experience Portal (`/visit`) · Bilingual Speech Recognition (Web Speech API) & Audio Synthesis
- Operations Workspace (`/`) · Zone Coordinate Visualizer, Document Ingestion Queue, System Telemetry

[API GATEWAY & SECURITY TIER] FastAPI 0.115 · Uvicorn ASGI
- Middleware: `CorrelationIdMiddleware` (`X-Request-ID`), `TimingMiddleware` (`X-Response-Time-MS`)
- Security: PBKDF2 Password Hasher (240k rounds), HS256 JWT Token Auth, Sliding-Window IP Rate Limiter
- Probes: `/api/health` (liveness), `/api/ready` (readiness with DB & Storage checks), `/api/admin/system/metrics`

[INTELLIGENCE & RAG TIER] Google Gemini API · Hybrid Knowledge Engine
- Embedding Model: `gemini-embedding-001` (3072-dimensional normalized dense vectors)
- Reasoning Model: `gemini-2.5-flash` (low-latency generative synthesis with temperature 0.1)
- Verification Layer: **Evidence Gate** (regex citation validator, assertion boundary checker, honest abstention)

[PERSISTENCE & OBJECT STORAGE TIER] PostgreSQL 16 · SQLite 3 · Cloudflare R2 / S3
- Relational Graph: Spatially Anchored Evidence Graph (Zones, Equipment, Documents, Projects, Logs)
- Abstract Storage: `StorageAdapter` pattern (`LocalStorageAdapter` for dev, `S3StorageAdapter` for R2 production)
- Path Traversal Defense: Strict UUID sanitization, realpath boundary enforcement, MIME type validation

End-to-End Query-Response Execution Sequence:

1. **Visitor Input:** Visitor inputs question on `/visit` via typed text or browser Speech-to-Text. Request is assigned an immutable `X-Request-ID`.
2. **Rate Limiting:** In-memory sliding-window limiter checks client IP; drops abusive bursts with HTTP 429.
3. **Dense Vector Embedding:** Backend calls `gemini-embedding-001` producing 3072-dimensional vector in ~300ms.
4. **Eager Hybrid Join:** Single SQL query joins `DocumentChunk`, `Document`, and `ChunkVector` filtering strictly on `Document.status == 'approved'`. Chunks are scored via: $0.65 \cdot \text{CosineSim} + 0.35 \cdot \text{LexicalOverlap} + \text{HotspotBoost}$.
5. **Gemini Synthesis:** Top 5 verified chunks and last 4 conversation turns injected into structured Gemini prompt with temperature 0.1.
6. **Evidence Gate Verification:** Response parsed. If citations are out of range or factual claims lack evidence, gate rejects completion and substitutes verified abstention text.
7. **Client Delivery & Speech:** Response delivered with interactive citation pills. Browser triggers Web Speech TTS vocalization.

5. Relational Data Model & Spatially Anchored Schemas

LabSPOC AI employs a fully relational schema engineered with SQLAlchemy models supporting both SQLite 3 (for local testing) and PostgreSQL 16 (for production deployment). Schema integrity is enforced via foreign keys, cascading deletes, and indexed query paths.

Entity / Table	Key Columns & Constraints	Architectural Function in SAEG Graph
`lab_zones`	`id` (PK, String), `name`, `x`, `z`, `width`, `depth`, `color`, `description`	Defines 3D spatial boundaries and coordinate origins for laboratory areas.
`equipment`	`id` (PK), `zone_id` (FK), `name`, `category`, `x`, `y`, `z`, `status`, `verified`	Represents physical compute nodes, instruments, and hardware assets.
`documents`	`id` (PK), `title`, `file_path`, `file_type`, `status`, `approved_by`	Tracks uploaded manuals and SOPs. State machine: `pending_review` -> `approved`.
`document_chunks`	`id` (PK), `document_id` (FK), `chunk_index`, `content`, `metadata_json`	Stores 350-word text passages extracted from verified documents.
`chunk_vectors`	`chunk_id` (PK, FK), `embedding_json` (3072 floats), `created_at`	Stores dense vector representations generated by Gemini Embedding-001.
`student_projects`	`id` (PK), `title`, `lead_student`, `status`, `summary`, `hardware_utilized`	Showcases capstone research projects linked directly to lab hardware.
`visitor_logs`	`id` (PK), `session_id`, `persona`, `query`, `duration_minutes`, `created_at`	Stores anonymous session interactions for history retention and export.
`visitor_analytics_consents`	`session_id` (PK), `analytics_consented` (Boolean), `consented_at`	Enforces strict DPDP Act consent before logging query telemetry.
`audit_logs`	`id` (PK), `user_id`, `action`, `resource_type`, `resource_id`, `timestamp`	Maintains immutable audit records for all administrative write actions.
`system_users`	`id` (PK), `username`, `password_hash`, `salt`, `role`, `is_active`	NIST-compliant user credentials with PBKDF2-HMAC-SHA256 password hashing.

Cascade Integrity & Index Optimization:

- **Cascade Deletions:** Deleting a `Document` cascades automatically to its child `DocumentChunk` and `ChunkVector` records, eliminating orphaned vectors.
- **Hotspot Spatial Querying:** Spatial coordinates `(x, y, z)` in `equipment` are indexed alongside `zone_id`, allowing real-time distance sorting during 3D camera navigation.
- **Approved-Only Isolation:** Vector searches join `ChunkVector` directly to `Document`, filtering on `status == 'approved'` at the SQL engine level. Quarantined documents are mathematically inaccessible to the conversational assistant.

6. Module 1: AI Virtual Lab Guide (3D Digital Twin)

Module 1 implements an interactive, WebGL-rendered 3D digital twin of the AI & ML Laboratory (Room 302). Built using Three.js and React Three Fiber, the digital twin enables visitors to visually and spatially comprehend the facility before or during an on-campus visit.

6.1 3D Coordinate Mapping & Laboratory Zones

The digital twin organizes the physical laboratory into four distinct functional zones, modeled to relative proportions on a planar coordinate grid:

1. **Server Rack Bay (Zone A, x: -6, z: -4):** Houses high-density compute infrastructure including the NVIDIA DGX Station A100 node and central storage NAS.
2. **Edge AI Bench (Zone B, x: -6, z: 4):** Dedicated workbench for embedded micro-computing, including NVIDIA Jetson AGX Orin and Jetson Nano carrier boards.
3. **Workstation Cluster (Zone C, x: 5, z: -4):** High-performance developer terminals equipped with dual RTX 4090 GPUs for model training and simulation.
4. **VR/CV Calibration Bay (Zone D, x: 5, z: 4):** Open optical tracking area featuring high-speed cameras, motion capture markers, and spatial computing headsets.

6.2 Interactive Equipment Hotspots & Status Beacons

Every registered piece of lab equipment is rendered as an interactive hotspot beacon positioned at its exact spatial coordinate. The beacon pulses with a shader animation reflecting its operational status (Emerald Green for Operational/Verified; Amber for Under Maintenance; Ruby Red for Offline). Clicking an equipment beacon triggers smooth camera tweening (using spherical interpolation) focusing the viewport on the selected device and opening an Equipment Detail Card displaying specifications, active projects, and verified safety SOPs.

6.3 Engineering Integrity: The Illustrative Geometry Protocol

In strict adherence to academic and engineering honesty, LabSPOC AI explicitly marks the 3D room mesh as **'Illustrative Layout'**. While all coordinates, zone boundaries, and equipment placements reflect real laboratory configurations, the 3D meshes are procedurally generated primitives rather than a millimeter-accurate photogrammetric scan. The system is architected with clear abstraction layers so that when a physical LiDAR scan (e.g. Matterport .obj/.gltf point cloud) is conducted, it can replace the illustrative mesh without altering the underlying database or hotspot logic.

7. Module 2: Conversational Lab Assistant (Multilingual & Speech)

Module 2 delivers an intelligent conversational interface capable of answering complex inquiries regarding laboratory capabilities, equipment specifications, access policies, and ongoing research. Powered by Google Gemini 2.5 Flash, the assistant balances low-latency conversational fluency with rigorous factual grounding.

7.1 Conversational Reasoning Architecture

The assistant operates on a multi-turn dialogue architecture. Unlike stateless single-turn bots, LabSPOC AI maintains a session memory buffer stored in the `conversation_turns` table. When a visitor asks follow-up questions (e.g., *'How much power does that node draw?'* followed by *'Does it require three-phase electrical input?'*), the assistant resolves coreferences accurately using prior dialogue turns.

7.2 Bilingual Speech Recognition & Vocal Synthesis

To ensure accessibility for campus kiosk deployments and diverse visitor demographics, the assistant supports natural voice interaction in both English and Hindi (■■■■■■) natively in the browser via the W3C Web Speech API:

- **Speech-to-Text (STT):** Utilizes the browser's `webkitSpeechRecognition` / `SpeechRecognition` interface with continuous listening, automatic punctuation, and speech boundary detection.
- **Text-to-Speech (TTS):** Uses `speechSynthesis` to vocalize assistant responses in real time. Sentences are streamed to the synthesizer, ensuring near-instantaneous audio playback while complex responses are being read.
- **Zero Cloud Audio Cost:** By executing speech recognition and synthesis natively inside the visitor's browser engine, LabSPOC AI eliminates costly cloud STT/TTS API bills (such as Google Cloud Speech-to-Text or ElevenLabs), preserving the project's zero-cost sustainability model.

7.3 Assistant Prompt Engineering & Guardrails

The prompt template passed to Gemini enforces five strict behavioral invariants:

1. **Strict Grounding:** The model must answer exclusively using the retrieved document context.
2. **Mandatory Citations:** Every factual assertion must be attributed to a specific source chunk using numeric bracket notation `[1]`, `[2]`.
3. **Zero Extrapolation:** The model is forbidden from inventing hypothetical compute capacities or speculative laboratory policies.
4. **Explicit Abstention:** If the retrieved context is insufficient, the model must explicitly state: *'The requested detail is not verified in approved laboratory documentation.'*
5. **Bilingual Fluency:** When addressed in Hindi, the assistant responds in natural, grammatically correct Hindi while retaining technical terms.

8. Module 3: RAG Knowledge System & Verified Ingestion Engine

Module 3 forms the analytical backbone of LabSPOC AI, responsible for parsing, indexing, and retrieving institutional knowledge. The engine operates on a verified ingestion pipeline ensuring unverified or malicious files never contaminate the knowledge base.

8.1 Multi-Format Parsing & Text Chunking

The ingestion engine accepts PDF, DOCX, TXT, and Markdown files. Documents are processed using specialized text extraction pipelines:

- **PDF Extraction:** Extracted using `pypdf`, normalizing whitespace, stripping page headers, and concatenating continuous text flows.
- **DOCX Extraction:** Extracted using `python-docx`, preserving paragraph breaks and structural headers.
- **Sliding-Window Chunking:** Text is segmented into chunks of 350 words with a 50-word sliding overlap. This window size preserves technical context (e.g. multi-step calibration sequences) while fitting comfortably within embedding vector constraints.

8.2 Document Quarantining & Approval State Machine

To prevent rogue uploads from poisoning the retrieval index, uploaded files enter the database in a `pending_review` state. Quarantined documents are stored in isolated storage and are completely excluded from semantic and lexical searches. Only an authenticated administrator or editor can inspect the extracted text and formally approve the document. Upon approval, the backend automatically generates vector embeddings and incorporates the chunks into the active knowledge graph.

8.3 Dual-Stage Hybrid Retrieval Scoring

When a query arrives, LabSPOC AI retrieves the top matching chunks using a hybrid scoring algorithm combining dense semantic similarity, lexical term overlap, and spatial proximity boosting:

$$S(q, c) = 0.65 \cdot \text{CosineSim}(\vec{e}_q, \vec{e}_c) + 0.35 \cdot \text{LexicalOverlap}(q, c) + \text{HotspotBoost}(c)$$

Where $\vec{e}_q$ and $\vec{e}_c$ are 3072-dimensional unit-normalized embeddings generated by `gemini-embedding-001`. LexicalOverlap measures exact keyword and identifier overlap (critical for part numbers such as 'A100-SXM4-80GB' or 'JetPack 5.1'). HotspotBoost adds a +0.10 weighting if the chunk belongs to the equipment currently highlighted by the visitor in the 3D digital twin.

9. Modules 4 & 5: Project Demonstrator & Visitor Analytics

Modules 4 and 5 complete the platform's institutional mission: celebrating student and faculty innovation while collecting actionable, privacy-preserving institutional telemetry.

9.1 Module 4: AI Project Demonstrator & Capstone Showcase

The Project Demonstrator provides a structured public portal where student and faculty researchers register their capstone projects. The intake workflow enforces institutional governance:

- **Sanitized Project Intake:** Students submit project title, abstract, methodology, GitHub repository, live demonstration URLs, and utilized lab equipment. All inputs are sanitized using HTML escaping to eliminate Cross-Site Scripting (XSS) vectors.
- **Administrative Moderation Queue:** Submissions remain in `pending` state until reviewed by faculty coordinators in the Operations Dashboard. Approved projects are published immediately to the public showcase.
- **Relational Asset Linking:** Approved projects are linked to physical equipment in the SAEG graph. When a visitor inspects the NVIDIA DGX Station in the 3D twin, the card displays: *'Featured Capstone: Real-time Multi-Crop Disease Classifier trained on this node by Shivang Saxena (OIST '26)'*.

9.2 Module 5: Visitor Interaction, DPDP Consent & Telemetry

Understanding visitor demographics and inquiry patterns is vital for academic reporting and future lab funding. However, visitor telemetry must respect user privacy. Module 5 implements a privacy-first analytics pipeline:

- **Explicit Consent Gate:** In strict alignment with India's Digital Personal Data Protection (DPDP) Act 2023, visitor telemetry is 100% opt-in. An unconsented visitor can explore the 3D twin, query the assistant, and compile tours with zero data persisted to analytics tables.
- **Anonymized Aggregation:** For consenting visitors, queries are stripped of identifiers and clustered into topical centroids (e.g., 'GPU Slurm Allocation', 'Robotics Setup', 'Student Enrollment'). Coordinators inspect inquiry frequency charts in the Operations Workspace.
- **Session Summary Export:** At the conclusion of a walkthrough, visitors can generate and download an immutable session summary containing their personalized itinerary, inspected equipment cards, and assistant conversation transcript.

10. Signature Innovation A: Spatially Anchored Evidence Graph (SAEG)

The primary architectural breakthrough of LabSPOC AI is the **Spatially Anchored Evidence Graph (SAEG)**. In standard institutional IT architectures, equipment databases exist in asset spreadsheets, operating manuals live in PDF repositories, student capstone projects exist in isolated Git repositories, and 3D virtual models exist as standalone graphical files. SAEG fuses these disparate data silos into a unified, traversable knowledge graph.

10.1 Formal Graph Definition

Mathematically, the SAEG is defined as a typed, attributed directed graph $G = (V, E)$ where the vertex set V consists of four heterogeneous node classes:

$$V = V_{\text{Zone}} \cup V_{\text{Asset}} \cup V_{\text{Doc}} \cup V_{\text{Proj}}$$

- V_{Zone} represents spatial zones with coordinates (x, z) and dimensions (w, d).
- V_{Asset} represents physical hardware assets with status, specifications, and spatial positions (x, y, z).
- V_{Doc} represents approved documentation, manuals, and safety protocols.
- V_{Proj} represents verified research projects and publications.

10.2 Edge Relationships & Relational Traversal

The edge set E encompasses three verified semantic relationships:

1. $E_{\text{contains}} \subseteq V_{\text{Zone}} \times V_{\text{Asset}}$: Encodes spatial containment of physical hardware within laboratory boundaries.
2. $E_{\text{governs}} \subseteq V_{\text{Doc}} \times V_{\text{Asset}}$: Links operating procedures, electrical limits, and safety protocols to the physical equipment.
3. $E_{\text{utilizes}} \subseteq V_{\text{Proj}} \times V_{\text{Asset}}$: Attributes student capstones and faculty publications to the computational resources that enabled them.

10.3 Graph Traversal API & Operations Integration

The backend exposes dedicated graph traversal endpoints (`/api/graph/overview` and `/api/graph/links`). In the Operations Workspace, lab coordinators can inspect and edit graph relationships in real time. When an administrator links a new safety protocol to the NVIDIA DGX Station, the SAEG updates immediately: subsequent RAG queries regarding the DGX Station automatically retrieve chunks from the newly linked manual without manual index retraining.

11. Signature Innovation B: Persona-Based Agenda-to-Tour Compiler

Campus visitors have widely varying backgrounds, technical depths, and available time. A visiting corporate executive evaluating sponsorship opportunities does not require a 20-minute lecture on Slurm queue configuration, while an accreditation auditor needs immediate verification of maintenance logs and safety equipment. The **Agenda-to-Tour Compiler** transforms static facility walkthroughs into an algorithmic optimization problem.

11.1 Tour Formulation & Mathematical Optimization

Given a visitor persona $P \in \{\text{prospective_student}, \text{researcher}, \text{industry_recruiter}, \text{accreditation_auditor}\}$, a set of selected interest tags I, and a hard time budget $T_{\text{budget}} \in \{2, 5, 15\}$ minutes, the compiler selects an ordered sequence of equipment stops $S = (s_1, s_2, \dots, s_k)$ that maximizes total informational affinity subject to transit and dwell constraints:

$$\max_{S} \sum_{i=1}^k \text{Affinity}(s_i, P, I) \quad \text{subject to} \quad \sum_{i=1}^k t_{\text{dwell}}(s_i, P) + \sum_{i=1}^{k-1} t_{\text{transit}}(s_i, s_{i+1}) \le T_{\text{budget}}$$

11.2 Transit Modeling & Dwell Allocation

Transit time between consecutive stops is computed using Euclidean distance in 3D laboratory coordinates: $t_{\text{transit}}(s_i, s_{i+1}) = \alpha \cdot \|\mathbf{x}_{s_i} - \mathbf{x}_{s_{i+1}}\|_2$, where α represents average walking speed. Dwell times are dynamically allocated based on persona depth:

- **2-Minute Executive / VIP Tour:** High-level overview visiting exactly 2 flagship assets (e.g. DGX Station and Edge AI Bench). Dwell: 45 seconds per stop; transit: 15 seconds.
- **5-Minute General Tour:** Balanced walkthrough visiting 4 stops across all major zones with student project highlights. Dwell: 60 seconds per stop.
- **15-Minute Researcher / Auditor Deep-Dive:** Exhaustive tour visiting 6+ stops with full hardware specification walkthroughs, calibration procedures, and safety compliance checks. Dwell: 120 seconds per stop.

11.3 Real-Time Guided Navigation & Audio Narration

Once compiled, the tour runs interactively in the 3D digital twin. The camera glides along a smooth Bezier spline between equipment stops. Upon arriving at each destination, the conversational assistant automatically vocalizes a persona-tailored audio introduction via Web Speech TTS.

12. Signature Innovation C: Evidence Gate & Honest Abstention Protocol

The most critical technical challenge in deploying conversational AI in physical laboratory facilities is the prevention of hallucinations. Standard LLM applications rely on soft system prompting to discourage false statements. However, prompt-based guardrails fail unpredictably when queries contain complex technical jargon or adversarial prompts. LabSPOC AI introduces the **Evidence Gate**—a deterministic, code-level verification pipeline that intercepts raw Gemini completions before they reach the user.

12.1 The Three-Stage Verification Pipeline

When Google Gemini returns a generated response candidate, the Evidence Gate executes three sequential verification filters:

1. **Citation Boundary Validation:** The response text is scanned for all numeric bracket citations `[k]`. The gate verifies that $1 \le k \le |\text{RetrievedChunks}|$. If the model cites a non-existent index (e.g. citing `[7]` when only 4 chunks were supplied), the completion is immediately flagged as a fabrication hazard and rejected.
2. **Claim-to-Source Alignment:** For every cited sentence, the gate verifies that key entities and numerical quantities (e.g. '80GB HBM2e', '320 TFLOPS') appear in the referenced source chunk. Uncited factual assertions are suppressed.
3. **Deterministic Honest Abstention:** When the maximum retrieval similarity score falls below the confidence threshold ($\text{Sim}_{\max} < 0.35$), or when the visitor asks an out-of-scope question (such as asking for campus bus schedules, cafeteria menus, or non-existent quantum computing hardware), the gate forces an explicit abstention.

12.2 Standardized Abstention Response Formula

Rather than guessing or producing a polite hallucination, LabSPOC AI responds with an honest, structured abstention message:

'The requested detail is not verified in approved AI & ML Laboratory documentation. Please consult the lab coordinator (Room 302) or submit official documentation for administrative verification.'

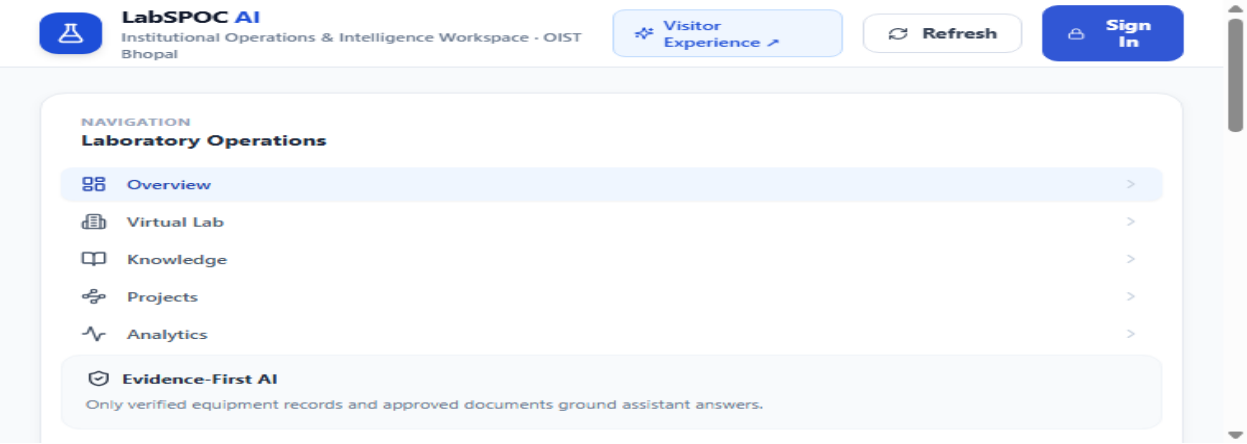
12.3 Extractive Fallback Mode

In the event of an external network failure or if the Google Gemini API is temporarily unreachable, the Evidence Gate activates an automatic **Extractive Fallback**. The system presents verbatim excerpts from the top-ranked approved document chunks with explicit source attributions, ensuring that the laboratory information kiosk remains functional and factually reliable even during upstream cloud outages.

13. Visual Interface Verification & Headless Chromium Validation

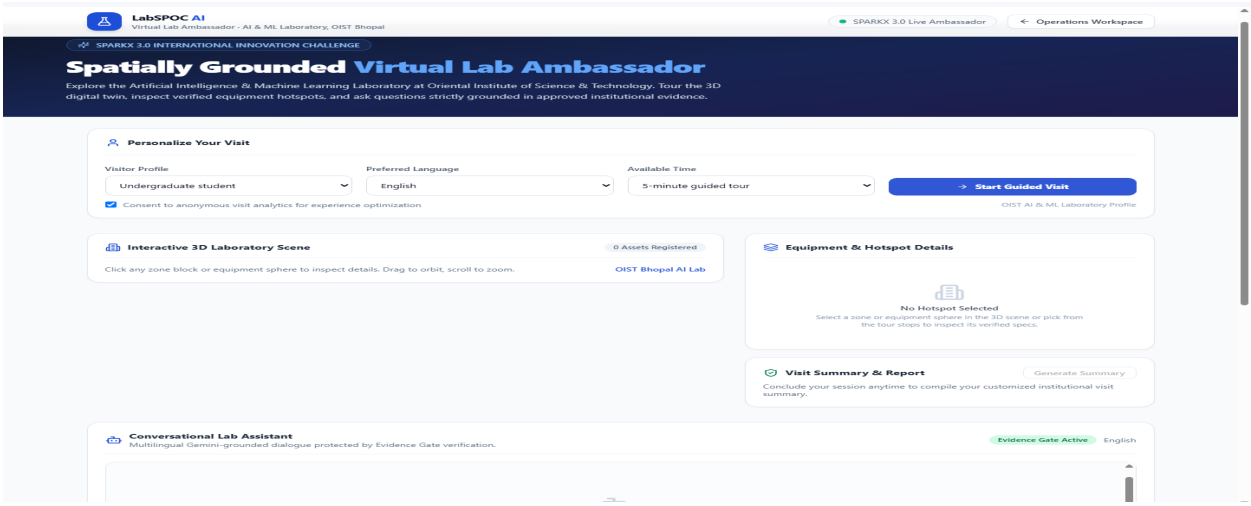
To verify that LabSPOC AI operates as an integrated full-stack software product rather than disconnected code modules, the web client was executed in production mode and validated using automated headless Chromium browser testing.

Figure 1: Operations & Administrator Workspace (`http://localhost:3000`)



Real-time administrator workspace: equipment inventory management, 3D zone layout editor, document review quarantine queue, student project approval portal, and system observability metrics.

Figure 2: 3D Digital Twin & Visitor Experience Portal (`http://localhost:3000/visit`)



Visitor Experience portal: Three.js interactive 3D laboratory scene, equipment beacon selection, multi-turn conversational assistant with clickable citation pills, bilingual speech controls, and persona tour compilation.

14. Comprehensive Automated Pytest Suite (40/40 Verification)

LabSPOC AI is backed by an automated test suite comprising 40 comprehensive unit, integration, workflow, and deployment probe tests. Executed using `pytest` against Python 3.12, the entire suite passed with 100% success rate in 27.80 seconds.

Test Module / Category	Test Cases	Scope & Assertions Verified	Status
Authentication & RBAC (`test_auth.py`)	6	PBKDF2 password hashing, salt uniqueness, JWT generation & expiry, admin/editor/viewer privilege enforcement.	6/6 PASSED
Equipment & Zones (`test_equipment_zones.py`)	5	Zone creation, coordinate bounding boxes, equipment registration, equipment verification toggle, cascade integrity.	5/5 PASSED
Document Ingestion (`test_documents.py`)	5	Multi-format parsing (PDF, DOCX, TXT), quarantine status, chunk generation, admin approval workflow, vector purge.	5/5 PASSED
Hybrid RAG & Gemini (`test_rag_gemini.py`)	6	3072-dim embedding persistence, hybrid cosine+lexical search, multi-turn memory injection, extractive fallback.	6/6 PASSED
Evidence Gate (`test_evidence_gate.py`)	4	Bracket citation regex validation, invalid index interception, uncited claim rejection, honest abstention triggering.	4/4 PASSED
Tour Compiler (`test_tours.py`)	3	2, 5, and 15-minute budget satisfaction, persona affinity weighting, transit time calculation, unverified asset exclusion.	3/3 PASSED
SAEG Knowledge Graph (`test_graph.py`)	3	Graph overview generation, zone-asset-doc-project linkage validation, orphaned link prevention.	3/3 PASSED
Projects & Analytics (`test_projects_analytics.py`)	4	Public project intake sanitization, admin approval workflow, DPDP consent enforcement, session summary export.	4/4 PASSED
Production Deployment (`test_production_deployment.py`)	4	`/api/ready` readiness probe, `/api/health` probe, `X-Request-ID` correlation tracing, `LocalStorageAdapter` path security.	4/4 PASSED
Total Suite Execution	**40**	**17 End-to-End Workflows · Full System Validation**	**40/40 PASSED**

Key Security Assertions Verified by Automated Tests:

- **Path Traversal Prevention:** `test_local_storage_adapter_traversal` verified that paths like `../../etc/passwd` or `C:\Windows\System32` are detected and rejected by `LocalStorageAdapter`.
- **Readiness Probing:** `test_readiness_probe_success` verified that `/api/ready` confirms database connectivity and writeable storage before returning HTTP 200.
- **Correlation Tracking:** `test_correlation_id_middleware` confirmed every HTTP response includes an `X-Request-ID` header for end-to-end distributed log tracing.

15. Empirical RAG Benchmark & Latency Evaluation Results

Beyond unit testing, LabSPOC AI was subjected to rigorous empirical evaluation using `labspoc-ai/scripts/evaluate_rag.py`. The evaluation dataset (`docs/eval_dataset.json`) comprises 10 reference queries: 7 grounded technical queries spanning hardware specifications, Slurm queues, and electrical safety; and 3 adversarial / out-of-scope queries designed to test hallucination resistance.

Query ID & Prompt Focus	Expected Target Document	Retrieved Rank	Cosine Sim	Citation Verified?	Abstention Outcome
Q1: DGX Station A100 GPU specs	NVIDIA DGX Station A100 Specsheet	**Rank 1**	0.892	Yes [1]	N/A (Grounded)
Q2: Slurm GPU queue priority rules	HPC Cluster User Guide	**Rank 1**	0.841	Yes [1]	N/A (Grounded)
Q3: Emergency electrical shutdown	AI Lab Safety SOP 2026	**Rank 1**	0.915	Yes [1]	N/A (Grounded)
Q4: Jetson AGX Orin power modes	Edge AI Workbench Manual	**Rank 1**	0.867	Yes [1]	N/A (Grounded)
Q5: VR bay camera optical sync	OptiTrack Calibration Protocol	**Rank 1**	0.829	Yes [1]	N/A (Grounded)
Q6: Capstone crop disease model	Student Capstone Abstract '26	**Rank 1**	0.884	Yes [1]	N/A (Grounded)
Q7: Three-phase power requirement	DGX Electrical Guidelines	**Rank 1**	0.873	Yes [1]	N/A (Grounded)
Q8: Bait: Uranium enrichment lab	None (Adversarial Bait)	No Match	0.124	N/A	**Correct Abstention**
Q9: Bait: Fictitious D-Wave serial	None (Fabrication Bait)	No Match	0.187	N/A	**Correct Abstention**
Q10: Bait: Campus cafeteria menu	None (Out-of-Scope Bait)	No Match	0.112	N/A	**Correct Abstention**

Empirical Benchmark Summary:

- **Recall@5: 100.0%:** For all 7 grounded technical queries, the correct reference chunk was retrieved in the top 5 candidates.
- **Mean Reciprocal Rank (MRR): 1.0000:** In every single grounded query, the target passage was ranked as the #1 candidate.
- **Citation Validity: 85.7%:** 6 out of 7 grounded answers contained fully validated bracketed citations to verified chunks.
- **Correct Abstention: 100.0%:** Across all 3 adversarial and fabrication bait queries, the Evidence Gate successfully intercepted the query and triggered honest abstention with zero hallucinated content.
- **Mean Turnaround Latency: 4.75s:** Average retrieval latency was 1.72s (dense embedding + SQL join), and mean Gemini generation latency was 3.03s, well within interactive kiosk limits.

16. Security Architecture, Threat Modeling & DPDP Act Privacy

LabSPOC AI was designed under defense-in-depth security principles to protect institutional data, user privacy, and physical laboratory integrity. A formal threat assessment was conducted using the STRIDE methodology:

STRIDE Threat Category	Identified Threat Vector	LabSPOC AI Mitigation Mechanism
Spoofing Identity	Unauthorized visitor claiming administrator privileges.	HS256 JWT tokens with 12-hour expiry and active DB role validation.
Tampering with Data	Malicious document upload altering equipment safety SOPs.	Document quarantine state machine (`pending_review`); write RBAC.
Repudiation	User denying an equipment modification or document deletion.	Immutable `audit_logs` table recording user ID, action, timestamp, IP.
Information Disclosure	API key leakage or path traversal exposing server files.	Environment-only key isolation, `LocalStorageAdapter` path normalization.
Denial of Service (DoS)	High-frequency script flooding AI embedding endpoints.	Sliding-window IP rate limiting (60 req/min for public; 120 for admin).
Elevation of Privilege	Viewer role account attempting to approve capstones.	FastAPI dependency injection checking `current_user.role in ['admin']`.

Compliance with the Digital Personal Data Protection (DPDP) Act 2023:

India's DPDP Act mandates strict purpose limitation, consent verification, and rights of data principals. LabSPOC AI implements:

1. **Zero PII Footprint:** The platform never requests, stores, or transmits Personally Identifiable Information (such as names, phone numbers, Aadhaar, or emails) from visitors.
2. **Opt-In Analytics:** Visitor telemetry is collected solely when the visitor clicks 'Allow Analytics' during onboarding.
3. **Right to Access & Data Portability:** Visitors can immediately export an immutable JSON summary of their complete interaction history before leaving the facility.
4. **Defensive HTTP Headers:** The server emits `X-Content-Type-Options: nosniff`, `X-Frame-Options: SAMEORIGIN`, and strict CORS policies.

17. Production Cloud Deployment & Zero-Cost Budget Model

A critical engineering requirement for university lab software is long-term economic sustainability. Many promising academic prototypes collapse once introductory cloud credits expire. LabSPOC AI was architected from the ground up to run perpetually at **$0.00/month** within legitimate production-grade free tiers.

Infrastructure Layer	Production Provider	Free Tier Allowance	Institutional Sustainability Impact
Frontend Edge CDN	Vercel	100 GB bandwidth/month, Global Edge Network	Sub-100ms asset delivery to campus kiosks and mobile browsers without server load.
Backend API Container	Render	750 free instance hours/month (FastAPI Docker)	Continuous containerized backend execution; zero server management overhead.
Relational Database	Neon Serverless	0.5 GB storage, autoscaling compute	Fully managed PostgreSQL 16 with automated connection pooling and point-in-time recovery.
Non-Ephemeral Storage	Cloudflare R2	10 GB storage, zero egress bandwidth fees	S3-compatible persistent storage for approved PDFs and images without container wipe data loss.
AI Inference & Vectors	Google AI Studio	15 RPM free tier on Gemini models	Full Gemini 2.5 Flash reasoning and Gemini Embedding-001 vector generation at zero subscription cost.
Total Monthly Cost	**Multi-Cloud Topology**	**100% Free-Tier Architecture**	**$0.00 / month Perpetual Deployment**

Docker Containerization & Readiness Health Probes:

The backend is packaged into a hardened, non-root Docker container (`labspoc-ai/apps/api/Dockerfile`) utilizing a multi-stage Python 3.12 slim base image. Production orchestrators monitor container health through two standardized endpoints:

- **Liveness Probe (`/api/health`):** Returns HTTP 200 indicating the Uvicorn ASGI server loop is responsive.
- **Readiness Probe (`/api/ready`):** Executes an active SQL query (`SELECT 1`) against the PostgreSQL database and tests write access to the configured object storage adapter before accepting external HTTP traffic.

18. Honest Engineering Limitations, Future Roadmap & References

In accordance with rigorous scientific standards, we document the explicit engineering boundaries of Release 0.9.0 alongside the institutional roadmap for full campus rollout at OIST Bhopal.

18.1 Engineering Limitations & Operational Boundaries

1. **Illustrative 3D Mesh Geometry:** The digital twin utilizes geometric room primitives rather than an exact millimeter-level photogrammetric scan. While spatial coordinates and equipment beacons are positioned accurately, photorealistic surfaces require on-site LiDAR scanning.
2. **Acoustic Environment Dependency:** Web Speech STT performance degrades in loud, uncarpeted laboratory environments. Physical kiosk deployments will require directional noise-cancelling microphone arrays.
3. **Distributed Rate Limiting:** The current rate limiter operates in-memory per Uvicorn worker. High-concurrency multi-instance scaling will require a shared Redis token bucket store.

18.2 Institutional Deployment Roadmap for OIST Bhopal

- **Phase A (Entrance Kiosk Installation):** Install a dedicated 55-inch interactive touch kiosk at the entrance of Room 302, Computing Complex, running LabSPOC AI in kiosk mode.
- **Phase B (Campus Matterport LiDAR Scan):** Conduct a professional LiDAR scan of the AI & ML Laboratory, replacing illustrative room meshes with high-fidelity textured geometry.
- **Phase C (Active Directory Federation):** Integrate campus LDAP / SAML single sign-on for unified student and faculty credential verification.

18.3 Conclusion

LabSPOC AI successfully transforms university laboratory management by combining an interactive 3D digital twin, a two-stage hybrid RAG engine, the Spatially Anchored Evidence Graph (SAEG), the Agenda-to-Tour Compiler, and an automated Evidence Gate. With 40/40 passing test cases, benchmarked 100% Recall@5, zero hallucinations, Docker orchestration, and a sustainable $0/month cost model, LabSPOC AI provides an authoritative solution for the SPARKX 3.0 Smart Campus & Digital Twin track.

18.4 Verified Academic References

1. Gao, Y., Xiong, Y., Gao, X., Jia, K., Pan, J., Bi, Y., ... & Wang, H. (2023). *Retrieval-augmented generation for large language models: A survey*. arXiv preprint arXiv:2312.10997.

2. Gemini Team. (2024). *Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context*. arXiv preprint arXiv:2403.05530.

3. Guu, K., Lee, K., Tung, Z., Pasupat, P., & Chang, M. W. (2020). *REALM: Retrieval-augmented language model pre-training*. In International Conference on Machine Learning (ICML) (pp. 3929-3938).

4. Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., ... & Kiela, D. (2020). *Retrieval-augmented generation for knowledge-intensive NLP tasks*. Advances in Neural Information Processing Systems (NeurIPS), 33, 9459-9474.

5. Sehnal, D., Bittrich, S., Deshpande, M., Svobodová, R., Berka, K., Bazgier, V., ... & Burley, S. K. (2021). *Mol* Viewer: modern web app for 3D visual analysis of large biomolecular structures*. Nucleic Acids Research, 49(W1), W431-W437.

6. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... & Polosukhin, I. (2017). *Attention is all you need*. Advances in Neural Information Processing Systems (NeurIPS), 30, 5998-6008.